

Kernel-level Characterization, Chatbot, Jetson Nano

GPU Util (%)

100

0

Zoom-in View

Medium SM Occupancy

Large Kernels

Token i

Medium SM Occupancy

Large Kernels

Token i+1

Latency of token generation (34ms/token)

